

## Assignment No. 3

```
In [3]: #Importing the Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [4]: df = pd.read_csv("Churn_Modelling.csv")
```

```
In [5]: #Preprocesssing
df.isnull()
```

Out[5]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns

```
In [6]: df.isnull().sum()
```

```
Out[6]: RowNumber      0
CustomerId    0
Surname      0
CreditScore  0
Geography    0
Gender       0
Age          0
Tenure       0
Balance      0
NumOfProducts 0
HasCrCard    0
IsActiveMember 0
EstimatedSalary 0
Exited      0
dtype: int64
```



```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   RowNumber             10000 non-null  int64  
 1   CustomerId            10000 non-null  int64  
 2   Surname               10000 non-null  object  
 3   CreditScore           10000 non-null  int64  
 4   Geography             10000 non-null  object  
 5   Gender                10000 non-null  object  
 6   Age                   10000 non-null  int64  
 7   Tenure                10000 non-null  int64  
 8   Balance               10000 non-null  float64 
 9   NumOfProducts         10000 non-null  int64  
10   HasCrCard             10000 non-null  int64  
11   IsActiveMember        10000 non-null  int64  
12   EstimatedSalary       10000 non-null  float64 
13   Exited                10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [8]: df.dtypes
```

```
Out[8]: RowNumber      int64
CustomerId    int64
Surname       object
CreditScore   int64
Geography     object
Gender        object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object
```

```
In [9]: df.columns
```

```
Out[9]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
              'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [10]: #Dropping the unnecessary columns
df= df.drop(['RowNumber', 'CustomerId', 'Surname'], axis =1)
```



```
In [11]: df.head()
```


```
Out[11]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
In [12]: def visualization(x, y, xlabel):  
    plt.figure(figsize=(10,5))  
    plt.hist([x, y], color=['red', 'green'], label = ['exit','not_exit'])  
    plt.xlabel(xlabel,fontsize=20)  
    plt.ylabel("No. of customer", fontsize=20)  
    plt.legend()
```


```
In [13]: df_churn_exited = df[df['Exited']==1]['Tenure']  
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
In [14]: visualization(df_churn_exited,df_churn_not_exited, "Tenure")
```



```
In [15]: df_churn_exited2 = df[df['Exited']==1]['Age']  
df_churn_not_exited2 = df[df['Exited']==0]['Age']  
visualization(df_churn_exited2,df_churn_not_exited2, "Age")
```





```
In [16]: X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
               'IsActiveMember', 'EstimatedSalary', 'Exited']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
In [17]: df = pd.concat([df, gender, states], axis=1)
```

```
In [18]: #Splitting the training and testing Dataset
df.head()
```

Out[18]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Male	Germany	Spain
0	619	France	Female	42	2	0.00	1	1	1	101348.88	1	0	0	0
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0	0	0	1
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1	0	0	0
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0	0	0	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0	0	0	1

```
In [19]: X= df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
               'IsActiveMember', 'EstimatedSalary', 'Male', 'Germany', 'Spain']]
y=df['Exited']
```

```
In [20]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

```
In [21]: #Normalizing the values with mean as 0 and Standard Deviation as 1
from sklearn.preprocessing import StandardScaler
sc =StandardScaler()
X_train =sc.fit_transform(X_train)
X_test =sc.transform(X_test)
```



In [22]: `x_train`

```
Out[22]: array([[ 0.38628108,  0.48664962,  1.04235774, ...,  0.90244757,
                  1.72155741, -0.57339125],
                [ 0.80193718,  0.67749795, -0.68454375, ...,  0.90244757,
                  1.72155741, -0.57339125],
                [ 0.81232858, -0.18131955, -0.68454375, ...,  0.90244757,
                  -0.58086939, -0.57339125],
                ...,
                [-0.48659671,  1.15461879,  1.04235774, ...,  0.90244757,
                  -0.58086939, -0.57339125],
                [-0.67364195,  0.48664962,  0.69697744, ...,  0.90244757,
                  1.72155741, -0.57339125],
                [ 1.84107741, -0.37216788, -1.72068464, ..., -1.10809762,
                  -0.58086939,  1.74400987]])
```

In [23]: `x_test`

```
Out[23]: array([[ -0.33072568,  0.29580128, -1.02992404, ..., -1.10809762,
                  -0.58086939, -0.57339125],
                [-0.39307409,  0.29580128,  1.73311833, ..., -1.10809762,
                  -0.58086939, -0.57339125],
                [ 1.30072449,  1.25004295,  1.04235774, ...,  0.90244757,
                  -0.58086939, -0.57339125],
                ...,
                [ 0.27197566, -0.75386455, -0.68454375, ...,  0.90244757,
                  -0.58086939, -0.57339125],
                [ 0.83311138,  0.00952878,  0.35159714, ..., -1.10809762,
                  -0.58086939,  1.74400987],
                [-0.73599037, -0.18131955,  0.00621685, ...,  0.90244757,
                  -0.58086939,  1.74400987]])
```

```
In [24]: from sklearn.neural_network import MLPClassifier
ann = MLPClassifier(hidden_layer_sizes=(100,100,100),random_state = 0,max_iter =100, activation ='relu')
ann.fit(X_train,y_train)
```

C:\ProgramData\Anaconda\lib\site-packages\sklearn\neural\_network\\_multilayer\_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.  
warnings.warn(

```
Out[24]: MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100, random_state=0)
```

```
In [25]: y_pred =ann.predict(X_test)
```

```
In [26]: y_pred
```

```
Out[26]: array([1, 0, 0, ..., 0, 1, 0], dtype=int64)
```

```
In [30]: from sklearn.metrics import ConfusionMatrixDisplay,accuracy_score, classification_report
```


```
In [31]: y_test.value_counts()
```

```
Out[31]: 0    2396
         1     604
         Name: Exited, dtype: int64
```



```
In [32]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[32]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2a7aa29cc70>
```



```
In [39]: print(accuracy_score(y_test,y_pred))|
```

```
0.809
```

```
In [40]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.88	0.88	2396
1	0.53	0.53	0.53	604
accuracy			0.81	3000
macro avg	0.70	0.70	0.70	3000
weighted avg	0.81	0.81	0.81	3000